

Assignment 2 _Group 6

Baochen Hu & Ning Zhang

Kafka word count:

Python:

1. Start Zookeeper

```
bin/zookeeper-server-start.sh config/zookeeper.properties
```

2. configuration

```
/config/server.properties
```

```
host.name = localhost
```

```
advertised.host.name = localhost
```

3. Start Kafka Broker

```
bin/kafka-server-start.sh config/server.properties
```

4. create topic

```
bin/kafka-topics.sh --create --zookeeper
```

```
localhost:2181 --replication-factor 1 -- partitions 1 --topic test
```

```
bin/kafka-topics.sh --list --zookeeper localhost:2181
```

5. Start a Producer

```
bin/kafka-console-producer.sh --broker-list localhost:9092 --topic test
```

6. Start a consumer

```
bin/kafka-console-consumer.sh --zookeeper localhost:2181 --topic test  
--from-beginning
```

7. download a reliable jar, and put it in /external/kafka-assembly

```
wget http://search.maven.org/remotecontent?filepath=org/apache/spark/  
spark-streaming-kafka-assembly\_2.10/1.4.0/spark-streaming-kafka-assembly\_2.10-1.4.0.jar
```

6.

```
bin/spark-submit --jars
```

```
external/kafka-assembly/spark-streaming-kafka-assembly_2.10-1.4.0.jar
```

```
examples/src/main/python/streaming/kafka_wordcount.py localhost:2181
```

```
test
```

Scala:

bin/run-example org.apache.spark.examples.streaming.KafkaWordCount lo

calhost:2181 test-consumer-group test 1

```
-----  
Time: 2015-08-08 04:36:14  
-----
```

```
(u'hello', 3)  
( )  
-----
```

Flume word count

1.

apt-get [http://apache.arvixe.com/flume/1.6.0/apache-flume-1.6.0-bin.tar.g](http://apache.arvixe.com/flume/1.6.0/apache-flume-1.6.0-bin.tar.gz)

[z](#)

2.tar -xvz apache-flume-1.6.0-bin.tar.gz

3.add config.properties file.

a1.sources = r1

a1.sinks = k1

a1.channels = c1

Describe/configure the source

a1.sources.r1.type = netcat

a1.sources.r1.bind = 127.0.0.1

```
a1.sources.r1.port = 44444
```

```
# Describe the sink
```

```
#a1.sinks = avroSink
```

```
a1.sinks.k1.type = avro
```

```
a1.sinks.k1.channel = memoryChannel
```

```
a1.sinks.k1.hostname = 127.0.0.1
```

```
a1.sinks.k1.port = 8989
```

```
#a1.sinks.k1.type = logger
```

```
# Use a channel which buffers events in memory
```

```
a1.channels.c1.type = memory
```

```
a1.channels.c1.capacity = 5
```

```
a1.channels.c1.transactionCapacity = 5
```

```
# Bind the source and sink to the channel
```

```
a1.sources.r1.channels = c1
```

a1.sinks.k1.channel = c1

4. bin/flume-ng agent --conf conf --conf-file

conf/flume-spark-integration.properties --name a1

-Dflume.root.logger=INFO,console

5. open another terminal: telnet localhost 44444

6. open another terminal: bin/run-example

org.apache.spark.examples.streaming.FlumeEventCount 127.0.0.1 8989

NOTE: This should be ran first.

```
-----  
Ti root@li796-54: ~/apache-flume-1.6.0-bin — ssh — 159x20  
-----
```

Received 1 flume events.

```
-----  
Time: 1438999076000 ms  
-----
```

Received 0 flume events.

```
-----  
Time: 1438999078000 ms  
-----
```

Received 0 flume events.

```
-----  
Time: 1438999080000 ms  
-----
```

Received 0 flume events.

```
Trying ::1...
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.
hello
OK
ni
OK
Connection closed by foreign host.
root@li796-54:~# telnet localhost 44444
Trying ::1...
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.
hello
OK
hhllo
OK
^CConnection closed by foreign host.
```

HDFS word count:

hadoop fs -ls /user

hadoop fs -mkdir /user/ningzhang/sparkStreaming

scala:

Counts words in new text files created in the given directory

* Usage: HdfsWordCount <directory>

* <directory> is the directory that Spark Streaming will use to find and read new text files.

*

* To run this on your local machine on directory `localdir`, run this example

* \$ bin/run-example \

* org.apache.spark.examples.streaming.HdfsWordCount localdir

*

* Then create a text file in `localdir` and the words in the file will get counted.

*/

bin/run-example org.apache.spark.examples.streaming.HdfsWordCount

hdfs://ec2-52-2-200-145.compute-1.amazonaws.com/data/

hdfs dfs -copyFromLocal /home/hadoop/txtFile/hello.txt

hdfs://ec2-52-2-200-145.compute-1.amazonaws.com/data/hello1

Time: 1438993764000 ms

(hi,3)
(hello,2)

MQTT word count:

Install mosquitto:

1.<http://mosquitto.org/download/>

For Centos

Add the CentOS mosquitto repository to YUM's list of repositories

\$ cd /etc/yum/yum.repos.d

```
$ sudo
```

```
wget http://download.opensuse.org/repositories/home:/oojah:/mqtt/Cent
```

```
OS_CentO...
```

```
$ sudo yum update
```

```
$ sudo yum install mosquitto
```

As of writing, no init.d script exists for the CentOS distribution of mosquitto.

However, it is a simple enough matter to set it running as a daemon, you'll just need to restart it yourself whenever your machine gets restarted

```
1.$ sudo su
```

```
2.$ /usr/sbin/mosquitto -d -c /etc/mosquitto/mosquitto.conf >
```

```
    /var/log/mosquitto.log 2>&1
```

```
$mosquitto
```

To run this example locally, you may run publisher as

```
bin/run-example
```

```
org.apache.spark.examples.streaming.MQTTPublisher tcp://localhost:1883
```

```
foo
```



```
bin/run-example
```

```
org.apache.spark.examples.streaming.MQTTWordCount tcp://localhost:18
```

```
83 foo
```

TwitterPopularTags:

```
scala:
```

```
1.
```

```
./bin/run-example
```

```
org.apache.spark.examples.streaming.TwitterPopularTags
```

```
fLFWZfIOEZbUhPjhtfYvtmycy
```

```
iPGu1R7NGTt8SQkRpmFC1OfSs3VOC6vJwhzBP7yhABYTnzNvde
```

```
2612760793-cSRGyHgdIByaMeQZojisB6HYCJaPw9rgYsokbqjL
```

```
CACXdrOf7XDnWOUrc94fdawucDs56Ylz7GZRtRpFRDnil
```

<http://spark.apache.org/docs/latest/quick-start.html>

```
1. install sbt
```

```
sudo yum update
```

```
curl https://bintray.com/sbt/rpm/rpm | sudo tee
```

```
/etc/yum.repos.d/bintray-sbt-rpm.repo
```

```
sudo yum install sbt
```

2. dependency: simple.sbt

```
name := "Simple Project"
```

```
version := "1.0"
```

```
scalaVersion := "2.10.4"
```

```
libraryDependencies ++= Seq(
```

```
  "org.apache.spark" %% "spark-core" % "1.4.1",
```

```
  "org.apache.spark" %% "spark-streaming" % "1.4.1",
```

```
  "org.apache.spark" %% "spark-streaming-twitter" % "1.4.1"
```

```
)
```

```
vi TwitterPopularTags.scala
```

```
/* StreamingExamples.setStreamingLogLevels() */
```

3. sbt package

4. go to twitter developer to get token

<https://apps.twitter.com/app/new>

Consumer Key (API Key) fLFWZfIOEZbUhPjhtfYvtmycy

Consumer Secret (API

Secret) iPGu1R7NGTt8SQkRpmFC1OfSs3VOC6vJwhzBP7yhABYTnzNvde

Access

Token 2612760793-cSRGyHgdlByaMeQZojisB6HYCJaPw9rgYsokbqjL

Access Token Secret

CACXdrOf7XDnWOUrc94fdawucDs56Ylz7GZRtRpFRDnil

5.

./../usr/lib/spark/bin/spark-submit

twitterPopularTags/target/scala-2.10/simple-project_2.10-1.0.jar

fLFWZfIOEZbUhPjhtfYvtmycy

iPGu1R7NGTt8SQkRpmFC1OfSs3VOC6vJwhzBP7yhABYTnzNvde

2612760793-cSRGyHgdlByaMeQZojisB6HYCJaPw9rgYsokbqjL

CACXdrOf7XDnWOUrc94fdawucDs56Ylz7GZRtRpFRDnil

```
Popular topics in last 60 seconds (50 total):
#AlDubWeBelongTogether (2 tweets)
#flycosta (1 tweets)
#mizzoudg (1 tweets)
#騒音トラブル
http://t.co/WnXFwef8GC (1 tweets)
#sistas (1 tweets)
#online_ (1 tweets)
#Internet (1 tweets)
#里親募集 (1 tweets)
#Follow (1 tweets)
#football (1 tweets)
15/08/08 05:01:24 WARN BlockManager: Block input-0-1439010084200 replicated to only 0 peer(s) instead of 1 peers

Popular topics in last 10 seconds (50 total):
#AlDubWeBelongTogether (2 tweets)
#flycosta (1 tweets)
#mizzoudg (1 tweets)
#騒音トラブル
http://t.co/WnXFwef8GC (1 tweets)
#sistas (1 tweets)
#online_ (1 tweets)
#Internet (1 tweets)
#里親募集 (1 tweets)
#Follow (1 tweets)
#football (1 tweets)
```

ZeroMQ word count:

Install ZeroMQ:

<https://www.digitalocean.com/community/tutorials/how-to-install-zeromq-from-source-on-a-centos-6-x64-vps>

wget <http://download.zeromq.org/zeromq-2.1.10.tar.gz>

2.0.10

bin/run-example org.apache.spark.examples.streaming.SimpleZeroMQPu

blisher tcp://127.0.1.1:1234 foo.bar

```
bin/run-example org.apache.spark.examples.streaming.ZeroMQWordCount tcp://127.0.1.1:1234 foo
```

```
15/08/07 23:55:58 WARN BlockManager: Block input-0-1438991757105 replicated to only 0 peer(s) instead of 1 peers
15/08/07 23:55:59 WARN BlockManager: Block input-0-1438991757106 replicated to only 0 peer(s) instead of 1 peers
-----
Time: 1438991760000 ms
-----
(count,2)
(words,2)
(may,2)
```

Kinesis Word Count:

<https://github.com/apache/spark/blob/master/extras/kinesis-asl/src/main/scala/org/apache/spark/examples/streaming/KinesisWordCountASL.scala>

Set up Kinesis stream (see earlier section) within AWS. Note the name of the Kinesis stream and the endpoint URL corresponding to the region where the stream was created.

<https://console.aws.amazon.com/kinesis/home?region=us-east-1#stream-detail:myKinesisStream>

Set up the environment variables `AWS_ACCESS_KEY_ID` and `AWS_SECRET_KEY` with your AWS credentials.

<http://docs.aws.amazon.com/AWSSdkDocsJava/latest/DeveloperGuide/credentials.html>

```
$ export AWS_ACCESS_KEY_ID=AKIAJVZXGXB72L4TZUUQ
```

```
$ export AWS_SECRET_KEY=vJCOq6o9XiRqKCas9zmiyp7fFGgqOMRIoTMSp
```

XLE

```
./bin/run-example
```

```
org.apache.spark.examples.streaming.KinesisWordCountASL
```

```
myKinesisStream https://kinesis.us-east-1.amazonaws.com
```

```
./bin/run-example org.apache.spark.examples.streaming.KinesisWordCou
```

```
ntProducerASL myKinesisStream https://kinesis.us-east-1.amazonaws.com
```

```
1000 10
```

```
^C[hadoop@ip-172-31-37-15 spark]$ ./bin/run-example org.apache.spark.examples.streaming.KinesisWordCountPro
Putting records onto stream myKinesisStream and endpoint https://kinesis.us-east-1.amazonaws.com at a rate o
Sent 1000 records
Sent 1000 records
Sent 1000 records
Sent 1000 records
Sent 1000 records
Totals
(0,4976)
(1,4957)
(2,4978)
(3,5010)
(4,4945)
(5,5107)
(6,5175)
(7,4949)
(8,4948)
(9,4955)
^C[hadoop@ip-172-31-37-15 spark]$
```

Time: 1439060860000 ms

(4,152)
(8,151)
(0,153)
(5,150)
(9,141)
(1,162)
(6,161)
(2,155)
(7,136)
(3,169)

15/08/08 10:07:40 INFO scheduler.JobScheduler: Finished job stream